

Operating System (CSC 3150)

Tutorial 2

YUANG CHEN

E-MAIL: YUANGCHEN@LINK.CUHK.EDU.CN


Target

In this tutorial, we will practice write system in kernel mode.

- Kernel Object
- Insert and Remove Kernel Module
- Create Kernel Thread
- Compile Kernel
- System call execution

Process

- User Mode
- Kernel Mode



Kernel Object

- A loadable kernel module (or LKM) is an object file that contains code to extend the running kernel, or so-called base kernel
- LKMs are typically used to add support for new hardware and/or file systems, or for adding system calls.
- Most current Unix-like systems support loadable kernel modules, although they might use a different name for them,
 - for example: kernel extension (kext) in MacOS

Kernel Object Compiling (Makefile)

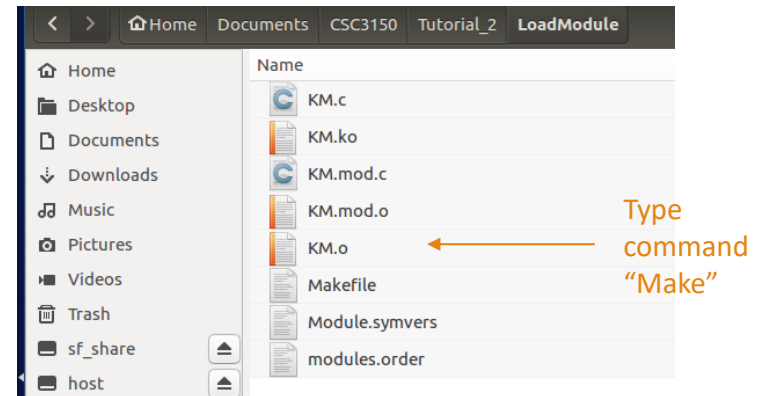
- Makefile

<http://www.cyberciti.biz/tips/compiling-linux-kernel-module.html>

- Build kernel object

```
(~/Documents/CSC3150/Tutorial_2/LoadModule) - gedit
Open [icon]
1 obj-m := KM.o
2 KVERSION := $(shell uname -r)
3 PWD := $(shell pwd)
4
5 all:
6     $(MAKE) -C /lib/modules/$(KVERSION)/build M=$(PWD) modules
7 clean:
8     $(MAKE) -C /lib/modules/$(KVERSION)/build M=$(PWD) clean
```

'KM' is name of the Kernel Object




If you type command "Make clean", it will clear all built files and leave original c file and makefile .

Insert and Remove Kernel Module

- Before insert the kernel object, you have to sign in the root account.
 - `$ sudo su`
- Insert module
 - `$insmod MODULE_NAME.ko`
- List the module you insert
 - `$lsmod`
 - `$lsmod | grep MODULE_NAME`
- Remember to remove your module
 - `$rmmod MODULE_NAME.ko`

Insert and Remove Kernel Module

```
Open ▾ 
1 #include <linux/init.h>
2 #include <linux/module.h>
3
4 MODULE_LICENSE("GPL");
5
6 static int KM_init(void) {
7     printk(KERN_INFO "Kernel Module initialization!\n");
8 }
9 return 0;
10 }
11
12 static void KM_exit(void) {
13     printk(KERN_INFO "Kernel Module exits!\n");
14 }
15 }
16
17 module_init(KM_init);
18 module_exit(KM_exit);
```

printk(): prints
the message into
kernel log

- Column 1: Module Name
- Column 2: Module Size
- Column 3: Used by
(denotes each module's
use count and a list of
referring modules)

```
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/LoadModule
[09/18/18]seed@VM:~/../LoadModule$ sudo su
[sudo] password for seed:
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule# insmod KM.ko
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule# lsmod | grep KM
KM                16384  0
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule# dmesg | tail -n 1
[ 5477.829462] Kernel Module initialization!
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule# rmmod KM.ko
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule# dmesg | tail -n 1
[ 5606.811308] Kernel Module exits!
root@VM:/home/seed/Documents/CSC3150/Tutorial_2/LoadModule#
```

grep: global
search regular
expression and
print out the line

dmesg: display
message buffer in
kernel

Kernel Thread

- Kthread creation:

- `struct task_struct *kthread_create(int (*threadfn)(void *data),
void *data,
const char *namefmt, ...);`
- The data argument will simply be passed to the thread function.
- The thread will not start running immediately. It will start to execute when returned task_struct is passed to wake_up_process().

- Kthread execution function:

- `int thread_function(void *data);`
- It can either call do_exit directly if it is a standalone thread for which no one will call kthread_stop()
- Or return when 'kthread_should_stop' is true (which means kthread_stop has been called).

Kernel Thread

- Return value:
 - It returns task_struct when executes successfully.
 - When fails, it returns ERR_PTR
- Kthread start execution with:
 - `int wake_up_process (struct task_struct * p);`
 - `ERR_PTR`
- A convenient function which creates and starts the thread:
 - `struct task_struct *kthread_run(int (*threadfn)(void *data),
 void *data,
 const char *namefmt, ...);`
 - Same as `kthread_create()` + `wake_up_process()`

Kernel Thread

Create a kernel thread to execute func

```
Open  [?]
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kthread.h>
MODULE_LICENSE("GPL");

static struct task_struct *task;

//implement test function
int func(void* data) {
    int time_count = 0;
    do {
        printk(KERN_INFO "thread_function: %d times", ++time_count);

    }while(!kthread_should_stop() && time_count<=30);
    return time_count;
}

static int __init KT_init(void){
    printk("KT module create kthread start\n");
    //create a kthread
    task=kthread_create(&func,NULL,"MyThread");

    //wake up new thread if ok
    if(!IS_ERR(task)){
        printk("Kthread starts\n");
        wake_up_process(task);
    }
    return 0;
}

static void __exit KT_exit(void){
    printk("KT module exits! \n");
}

module_init(KT_init);
module_exit(KT_exit);
```

GPL: General Public License.
Loading a proprietary or non-GPL-compatible LKM will set a 'taint' flag in the running kernel

```
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread# insmod KT.ko
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread# lsmod | grep KT
KT                16384  0
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread# rmmod KT.o
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread#
```

Kernel Thread

```
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread# clear all
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread# dmesg | tail -n 34
[37933.573361] KT module create kthread start
[37933.573796] Kthread starts
[37933.574623] thread_function: 1 times
[37933.574625] thread_function: 2 times
[37933.574625] thread_function: 3 times
[37933.574625] thread_function: 4 times
[37933.574626] thread_function: 5 times
[37933.574626] thread_function: 6 times
[37933.574626] thread_function: 7 times
[37933.574627] thread_function: 8 times
[37933.574627] thread_function: 9 times
[37933.574627] thread_function: 10 times
[37933.574628] thread_function: 11 times
[37933.574628] thread_function: 12 times
[37933.574629] thread_function: 13 times
[37933.574629] thread_function: 14 times
[37933.574629] thread_function: 15 times
[37933.574630] thread_function: 16 times
[37933.574630] thread_function: 17 times
[37933.574630] thread_function: 18 times
[37933.574631] thread_function: 19 times
[37933.574631] thread_function: 20 times
[37933.574631] thread_function: 21 times
[37933.574632] thread_function: 22 times
[37933.574632] thread_function: 23 times
[37933.574632] thread_function: 24 times
[37933.574633] thread_function: 25 times
[37933.574633] thread_function: 26 times
[37933.574633] thread_function: 27 times
[37933.574634] thread_function: 28 times
[37933.574634] thread_function: 29 times
[37933.574634] thread_function: 30 times
[37933.574635] thread_function: 31 times
[37948.414574] KT module exits!
root@VM: /home/seed/Documents/CSC3150/Tutorial_2/KernalThread#
```

Compile Kernel

- Download source code from
 - <http://www.kernel.org>
 - `$sudo apt-get install linux-source` (Or you could type in this command to download source code directly)

- Extract the source file to /home/seed/work
 - `cp KENEL_FILE.tar.xz /home/seed/work`
 - `cd /home/seed/work`
 - `$sudo tar xvf KENEL_FILE.tar.xz`

- Login root account and go to kernel source directory
 - `$sudo su`
 - `$cd /home/seed/work /KENEL_FILE`

Compile Kernel

- Clean previous setting and start configuration

- \$make mrproper
- \$make clean
- \$make menuconfig
- save the config and exit


configuration written to .config

- Build kernel Image and modules

- \$make bzImage
- \$make modules
- \$make -j NUM_CORE

Kernel: arch/x86/boot/bzImage is ready (#1)
root@VM:/usr/src/linux-4.10.14#

(you could use this command to replace above two commands)



Remark: Error in menuconfig

- Command “make menuconfig” does not working
 - Use command “**sudo apt-get install libncurses5-dev libssl-dev**” to install the tool

```
scripts/Makefile.host:124: recipe for target 'scripts/kconfig/mconf.o' failed
make[1]: *** [scripts/kconfig/mconf.o] Error 1
Makefile:546: recipe for target 'menuconfig' failed
make: *** [menuconfig] Error 2
root@VM:/usr/src/linux-4.10.14#
```

Compile Kernel

- Install kernel modules

- \$make modules_install

DEPMOD 4.10.14
root@VM:/home/seed/sdb4/linux-4.10.14#

- Install kernel

- \$make install

done
root@VM:/home/seed/sdb4/linux-4.10.14#

- Reboot to load new kernel

- \$reboot

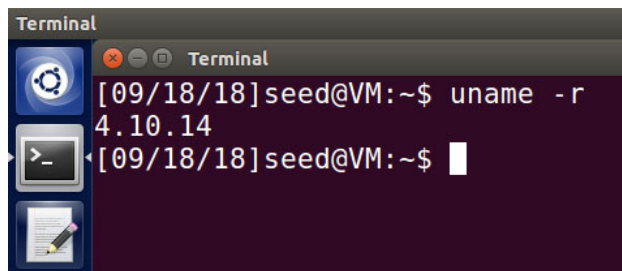
(When rebooting, you should select the updated kernel)

Compile Kernel

- Check exiting kernel version
 - \$uname -r




```
Terminal
[09/17/18]seed@VM:~$ uname -r
4.8.0-36-generic
[09/17/18]seed@VM:~$
```



```
Terminal
[09/18/18]seed@VM:~$ uname -r
4.10.14
[09/18/18]seed@VM:~$
```


System call execution (fork)



System call execution (fork)

- Calls `dup_task_struct()`, which creates a new kernel stack, `thread_info` structure, and `task_struct` for the new process.
- Calls `get_pid()` to assign an available PID to the new task.
- `copy_process()` then either duplicates or shares open files, filesystem information, signal handlers, process address space, and namespace.
- For more details
 - <https://elixir.bootlin.com/linux/v4.10.10/source/kernel/fork.c> (`do_fork`)

Export Symbol

- `EXPORT_SYMBOL()` helps you provide APIs to other modules/code.
- The functions which you EXPORT are available to the other modules/code.
- Your module will not load if the it's expecting a symbol(variable/function) and it's not present in the kernel.

References

- Loadable module kernel
 - https://en.wikipedia.org/wiki/Loadable_kernel_module

- Kthread_create()
 - <https://www.fsl.cs.sunysb.edu/kernel-api/re69.html>

- Linux commands
 - <http://www.runoob.com/linux/linux-command-manual.html> (Chinese)

References

- Compile kernel
 - <https://www.cnblogs.com/acm-icpcer/p/8029656.html> (version: Linux-4.10.14, Chinese)
 - <https://www.linux.com/learn/intro-to-linux/2018/4/how-compile-linux-kernel-0> (English)
 - http://www.berkes.ca/guides/linux_kernel.html (English)

- Extend storage in Virtual Box
 - <https://jingyan.baidu.com/article/d45ad148a1fab869542b8073.html> (Chinese)
 - <http://derekmolloy.ie/resize-a-virtualbox-disk/> (English)

Thank you

